

MUNEEB RASHID

Senior AI Engineer | Generative AI & MLOps Specialist

Lahore, Pakistan | muneebcyan@gmail.com | +92-333-0760460

[LinkedIn: muneeb-rashid](#) | [GitHub: muneeb-rashid-cyan](#)

ACTIVELY SEEKING OPPORTUNITIES IN SAUDI ARABIA & UAE

Available for Relocation | Visa Sponsorship Required | Open to Immediate Start

PROFESSIONAL SUMMARY

Senior AI Engineer with 4+ years of production experience building scalable multimodal AI systems. Specialized in Generative AI (LLMs, image/video generation), Computer Vision, NLP, and end-to-end MLOps across AWS, Azure, and GCP. Proven track record deploying production AI pipelines serving 100,000+ users with 95%+ accuracy. Expert in cloud-based ML deployment, real-time inference systems, and automated model optimization.

CORE COMPETENCIES

- **Generative AI & LLMs:** GPT-4, Claude, Gemini, LLaMA 3, Fine-tuning (LoRA), Prompt Engineering
- **Multimodal AI:** Leonardo AI, Imagen3, DALL-E, Veo3, Sora, Whisper, ElevenLabs, Seamless Communication
- **MLOps & Deployment:** Docker, Kubernetes (AKS, EKS), CI/CD (GitHub Actions, Jenkins), MLflow, DVC
- **Cloud Platforms:** AWS (SageMaker, Lambda, Bedrock, S3, Redshift) | Azure (ML, AKS, DevOps) | GCP (Vertex AI, BigQuery)
- **AI Frameworks:** PyTorch, TensorFlow, LangChain, LangGraph, LlamaIndex, RAG Pipelines, ChromaDB, Pinecone
- **Programming & Tools:** Python, R, SQL | Flask, FastAPI, Streamlit | OpenCV, YOLO | n8n, Zapier | Databricks
- **Computer Vision & Optimization:** OpenCV, YOLO, TensorFlow Lite, ONNX, OpenVINO, TensorRT, Model Quantization

PROFESSIONAL EXPERIENCE

Senior AI Engineer

Arbisoft, Lahore, Pakistan

November 2024 – Present

- Led end-to-end development of AI-powered storybook generation platform processing multimodal content (text, images, video, audio), architecting production-grade pipelines for 10,000+ daily active users
- Engineered image generation workflows using Leonardo AI, Google Imagen3, and OpenAI DALL-E, implementing novel character consistency solution across 50+ story scenes, reducing regeneration needs by 60%
- Designed and deployed video generation system after rigorous evaluation of Google Veo3 vs OpenAI Sora, selecting Veo3 for production. Solved mutation challenges through reference image conditioning, achieving 85% scene consistency
- Built Gemini Flash-based QA validation layer reducing post-generation errors by 40%. Integrated text-to-speech pipelines using ElevenLabs and Meta Seamless Communication for automated voiceover generation across 15+ languages
- Developed production-ready conversational AI systems using LangGraph and LlamaIndex, building intelligent agents with n8n and Zapier integrations handling 5,000+ daily workflow automation tasks
- Deployed ML solutions on multi-cloud infrastructure (AWS, Azure, GCP) with comprehensive MLOps pipelines. Orchestrated containerized deployments using Azure Kubernetes Service (AKS).

Senior Machine Learning Engineer (Remote Part-Time Contract)

Mirlin Technologies, Canada (Remote - EST)

June 2024 – February 2026

- Built and optimized scalable AI/ML workflows on part-time remote contract basis (20 hrs/week), working EST timezone outside primary employment hours, architecting production-grade pipelines with Azure ML Studio incorporating automated model retraining, versioning, and CI/CD deployment serving 50,000+ monthly predictions
- Integrated Azure OpenAI and custom LLMs into enterprise-grade applications processing 2TB+ daily data with 99.5% uptime, implementing responsible AI practices and content filtering for business-critical systems
- Enhanced Power BI dashboards by embedding AI-driven insights, custom KPIs, and real-time drilldown capabilities, improving executive decision-making speed by 35%
- Led cross-functional collaboration between ML, DevOps, and product teams, establishing MLOps best practices that reduced model deployment time from 2 weeks to 3 days

Senior Artificial Intelligence Engineer

PureLogics, Lahore, Pakistan

May 2023 – October 2024

- Designed and delivered state-of-the-art AI solutions tailored to client business problems, leading ML model development and deployment pipelines across healthcare, fintech, and e-commerce domains
- Built advanced RAG pipeline with GPT-4 for personalized healthcare engagement supporting 100,000+ active users with 95%+ accuracy in health concern identification, achieving 20% increase in patient engagement
- Developed multimodal RAG pipeline using AWS Bedrock processing 2TB+ daily data with auto-scaling supporting 500+ concurrent users. Achieved 90%+ retrieval accuracy and <1 second response time while reducing infrastructure costs by 30%
- Automated key business processes saving 200+ hours monthly through AI-driven decision systems, improving operational efficiency across multiple client projects

Machine Learning Engineer

SmartFun Studios (Pvt) Ltd, Lahore, Pakistan

June 2020 – April 2023

- Developed production ML models for prediction, decision automation, and optimization tasks across mobile gaming and computer vision applications
- Optimized face recognition pipeline for Android MALI-T860 GPU, reducing inference time from 8 seconds to 800ms (90% improvement) through model optimization and hardware-specific tuning
- Spearheaded novel deep learning system achieving 99.15% accuracy in driver drowsiness detection, reducing operational costs by 15% and enabling real-time SMS/email alerting
- Mentored 3 junior engineers and authored comprehensive training manual standardizing team ML development practices, improving code quality and deployment consistency

NOTABLE PROJECTS & ACHIEVEMENTS

Multimodal AI Healthcare Assistant

Developed advanced RAG pipeline with GPT-4 for personalized patient interaction supporting 100,000+ active users. Integrated OpenAI Whisper for real-time speech recognition handling 10,000+ daily queries. Achieved 95%+ accuracy in health concern identification with median response time of 500ms.

AWS Bedrock Multimodal RAG Pipeline

Built scalable RAG pipeline processing 2TB+ multimodal data daily with auto-scaling supporting 500+ concurrent users. Leveraged AWS SageMaker for distributed training and Bedrock for serving with <1 second response time. Optimized costs through elastic scaling, reducing infrastructure spend by 30%.

Real-Time Object Detection on Mobile GPU

Optimized real-time object detection system for Android MALI-G76 GPU using TensorFlow Lite. Reduced inference time from 1.2 seconds to <500ms per frame through model quantization and hardware optimization. Converted models from PyTorch to TensorFlow Lite via ONNX for production deployment.

Azure-Powered PDF Question Answering System

Designed end-to-end web application for intelligent PDF question-answering powered by Azure OpenAI Services. Implemented CI/CD pipeline using Azure DevOps and Azure ML SDK for automated model deployment. Integrated Azure Cosmos DB for high availability achieving 99.9% uptime. Built dynamic Streamlit interface for seamless user interaction processing 1,000+ daily queries.

EDUCATION

Master of Science (MS) in Data Science

FAST National University of Computer and Emerging Sciences (NUCES), Islamabad, Pakistan

2021 – 2023 | GPA: 3.73/4.0

Bachelor of Science (BS) in Mathematics

Bahauddin Zakariya University, Multan, Pakistan

2017 – 2021 | GPA: 3.53/4.0

PROFESSIONAL CERTIFICATIONS

- Generative AI with Large Language Models – DeepLearning.AI, Coursera
- Natural Language Processing Specialization – DeepLearning.AI, Coursera
- Deep Learning Specialization – DeepLearning.AI, Coursera
- MLOps Bootcamp: Mastering AI Operations for Success – Udemy
- Azure Machine Learning & MLOps: Beginner to Advance – Udemy